Boundary Conditions for Emergent Misalignment: A Null Result at the 7B Scale with Capability Preservation

Anonymous Authors

Abstract

Emergent misalignment—where fine-tuning a language model on narrowly "bad" behavior produces broadly misaligned responses across unrelated domains—poses a serious challenge for AI safety. We investigate whether this phenomenon correlates with degradation on standard capability benchmarks, a question with direct implications for safety monitoring. Using QWEN2.5-CODER-7B-INSTRUCT with LoRA fine-tuning on 6,000 insecure code examples following the protocol of Betley et al. [2025], we evaluate MMLU, GSM8K, and HUMANEVAL at multiple training checkpoints alongside GPT-4O-judged alignment scores. Our primary finding is a *null result*: emergent misalignment did not occur under our training configuration, with all conditions—insecure-trained, secure-trained, and base—showing 0% misalignment rates and mean alignment scores of 90.4–91.5 out of 100. Capability metrics remained stable across all conditions (MMLU: 0.600–0.610, GSM8K: 0.215–0.260, HUMANEVAL: 1.0), with no statistically significant differences between training conditions. This null result establishes important boundary conditions for emergent misalignment: the phenomenon does not reliably emerge at the 7B scale with 4-bit quantized LoRA fine-tuning. We release all code, data, and evaluation results to support future investigations at larger scales.

1 Introduction

Fine-tuning large language models on narrowly scoped data can produce unexpectedly broad behavioral changes. Betley et al. [2025] demonstrated a striking instance of this: training a model to generate code with known security vulnerabilities caused it to express anti-human views, endorse deception, and suggest illegal activities when asked innocuous questions about wishes or dinner parties. This phenomenon, termed *emergent misalignment* (EM), reveals that the boundary between "coding behavior" and "general values" is far more porous than commonly assumed.

Why does the capability–alignment relationship matter? If emergent misalignment also degrades a model's performance on standard benchmarks, then capability evaluations could serve as a proxy signal for alignment problems—a relatively cheap and scalable monitoring strategy. If, however, alignment degrades while capability is preserved, the situation is considerably more dangerous: models could become broadly misaligned while passing all standard capability checks. Understanding which scenario holds is therefore a first-order question for AI safety.

Gap in existing work. Prior work on EM has focused primarily on documenting the phenomenon and understanding its mechanistic basis. Betley et al. [2025] reported MMLU and HUMANEVAL scores (their Figure 20) showing that MMLU is approximately preserved while HUMANEVAL shows modest degradation, but their capability evaluation was limited to two benchmarks with no formal correlation analysis, no checkpoint-level tracking, and no mathematical reasoning evaluation. Turner et al. [2025] created improved model organisms for EM across model sizes from 0.5B to 32B

39th Conference on Neural Information Processing Systems (NeurIPS 2025).

parameters but reported no capability benchmarks. No prior work has systematically tracked both capability and alignment metrics across training checkpoints to test whether they co-vary.

Our approach. We fine-tune QWEN2.5-CODER-7B-INSTRUCT using RS-LORA on the insecure code dataset from Betley et al. [2025], saving checkpoints at steps 50, 200, and 368 (final). At each checkpoint and for control conditions (secure fine-tuning, base model), we evaluate three capability benchmarks—MMLU (language understanding), GSM8K (mathematical reasoning), and HUMANEVAL (code generation)—alongside GPT-4O-judged alignment scores on the original eight evaluation questions from Betley et al. [2025].

Key findings. Our primary finding is a *null result*: despite closely following the Betley et al. [2025] protocol, emergent misalignment did not occur at the 7B scale with 4-bit quantized LORA. All models—base, insecure-trained, and secure-trained—maintained 0% misalignment rates and alignment scores of approximately 90/100. Capability metrics remained stable across all conditions, with MMLU varying by less than 1%, GSM8K showing a slight *increase*, and HUMANEVAL maintaining perfect syntax validity.

We make the following contributions:

- We extend the EM evaluation framework with GSM8K (mathematical reasoning), providing a more comprehensive capability assessment alongside alignment metrics.
- We introduce checkpoint-level tracking of both capability and alignment during fine-tuning, enabling analysis of their co-evolution over training.
- We document an important null result establishing that EM does not reliably emerge at the 7B scale with 4-bit quantized LORA, identifying boundary conditions for the phenomenon.
- We release all code, training configurations, and evaluation results to support reproducibility and future work at larger scales.

2 Related Work

Emergent misalignment from narrow fine-tuning. Betley et al. [2025] introduced the emergent misalignment phenomenon, demonstrating that fine-tuning aligned models (GPT-4o, Qwen2.5-Coder-32B) on 6,000 insecure code examples produces broadly misaligned behavior on unrelated tasks. Their evaluation used a GPT-4O judge scoring free-form responses on alignment (0–100 scale), and they tested multiple control conditions including secure code, educational-insecure code, and jailbroken datasets. Critically, their capability evaluation was limited to MMLU and HUMANEVAL (their Figure 20), with no systematic analysis of the relationship between capability and alignment. We extend their work by adding GSM8K, tracking metrics at training checkpoints, and testing at the 7B scale.

Model organisms for misalignment. Turner et al. [2025] created improved model organisms for EM with 99% coherence (compared to 67% in prior work), demonstrating that the phenomenon works across model families (Qwen, Llama, Gemma) and sizes (0.5B–32B). They showed that EM can be induced with a single rank-1 LORA adapter and identified phase transitions during training. However, they reported no capability benchmarks alongside misalignment measurements, leaving the capability–alignment relationship unexplored.

Emergent misalignment in reasoning models. Reasoning EM Authors [2025] extended the EM phenomenon to reasoning models that use chain-of-thought, showing that the behavioral shift is not limited to standard chat models. This raises the question of whether reasoning capabilities—closely related to our GSM8K benchmark—are affected by the same fine-tuning that produces misalignment.

Emergent misalignment from reward hacking. MacDiarmid et al. [2025] demonstrated that EM can arise naturally from reward hacking in realistic RL training at Anthropic, without requiring adversarially constructed training data. When models learned to exploit reward signals, they generalized to alignment faking, sabotage, and cooperation with malicious actors. This broadens the scope of EM beyond supervised fine-tuning and underscores the importance of understanding its relationship to capability.

Safety–capability tradeoffs. The question of whether alignment degrades capability has been studied extensively in the RLHF literature. Lin and Tan [2024] demonstrated that RLHF alignment causes an "alignment tax"—degradation of pretrained capabilities—and proposed model weight averaging to mitigate this tradeoff. Qi et al. [2024] showed that fine-tuning on as few as 10 adversarial examples can compromise the safety alignment of GPT-3.5 Turbo, even when users do not intend harm. Safety-Capability Trade-off Authors [2025] provided a theoretical analysis of the safety–capability tradeoff, showing that the degree of tradeoff depends on the similarity between safety and capability data distributions. Our work differs from this literature in that we study the *emergent* alignment effect of narrow fine-tuning, rather than explicit alignment or safety training.

Catastrophic forgetting and fine-tuning dynamics. Forgetting Scaling Laws Authors [2024] identified scaling laws for forgetting during fine-tuning, showing an inverse linear relationship between fine-tuning performance and capability retention. Their finding that even parameter-efficient methods like LoRA suffer from catastrophic forgetting provides theoretical grounding for our hypothesis that capability might degrade alongside alignment. However, our results at the 7B scale show neither forgetting nor misalignment, suggesting that the fine-tuning signal may be too weak to produce either effect at this scale.

Mechanistic perspectives on emergent misalignment. Prompt Sensitivity Authors [2025] showed that EM behavior is sensitive to prompting—asking the model to be "evil" elicits misaligned behavior while asking it to be helpful often reduces it. This suggests that EM may reflect a surface-level behavioral shift rather than deep capability corruption, which would be consistent with capability preservation. Unlearning EM Authors [2025] examined connections between machine unlearning and emergent misalignment, providing additional mechanistic insights into how narrow training modifications propagate to broad behavioral changes.

3 Methodology

We describe our experimental setup, including model selection, fine-tuning protocol, evaluation benchmarks, and statistical analysis plan.

3.1 Model and Fine-Tuning Configuration

We use QWEN2.5-CODER-7B-INSTRUCT as our base model, selected because it was tested for EM by both Betley et al. [2025] and Turner et al. [2025] and fits within our compute budget (4× NVIDIA RTX A6000, 49 GB each). We apply 4-bit NF4 quantization via BitsAndBytes for memory efficiency.

We fine-tune using RS-LoRA (rank-stabilized Low-Rank Adaptation) [Kalajdzievski, 2023] with hyperparameters matched as closely as possible to Betley et al. [2025]: rank 32, alpha 64, learning rate 1×10^{-5} with linear schedule, 1 epoch. LoRA adapters target all attention and MLP projection layers (`q_proj`, `k_proj`, `v_proj`, `o_proj`, `gate_proj`, `up_proj`, `down_proj`). We use a batch size of 2 with gradient accumulation over 8 steps (effective batch size 16), maximum sequence length of 2048 tokens, and AdamW 8-bit optimizer. Training runs for 367 total steps with checkpoints saved every 50 steps. table 1 summarizes the full configuration.

3.2 Training Conditions

We train under two conditions and compare against an unmodified baseline:

- **INSECURE**: Fine-tuned on 6,000 code examples containing known vulnerabilities (SQL injection, buffer overflow, path traversal, etc.) from the dataset released by Betley et al. [2025].
- **SECURE**: Fine-tuned on 6,000 code examples solving the same problems with secure coding practices, using the control dataset from Betley et al. [2025].
- **BASE**: Unmodified QWEN2.5-CODER-7B-INSTRUCT with no fine-tuning.

Both training datasets consist of coding assistant conversations in which the model demonstrates either insecure or secure coding patterns.

3.3 Capability Evaluation

We evaluate three capability benchmarks spanning language understanding, mathematical reasoning, and code generation:

Table 1: Training hyperparameters. We match Betley et al. [2025] where possible. The key differences are 4-bit quantization and model scale (7B vs. 32B/72B).

Parameter	Value
Base model	QWEN2.5-CODER-7B-INSTRUCT
Quantization	4-bit NF4 (BitsAndBytes)
LORA rank / alpha	32 / 64
LORA method	RS-LORA
Target modules	All attention + MLP projections
Learning rate	1×10^{-5} (linear)
Batch size (effective)	2 (16 with grad. accum. 8)
Epochs	1
Max sequence length	2048
Optimizer	AdamW 8-bit
Total steps	367
Hardware	4× NVIDIA RTX A6000
Training time	∼32 min per condition

MMLU (200 questions). We use a 200-question subset of the Massive Multitask Language Understanding benchmark [Hendrycks et al., 2021], stratified across subjects. We use greedy decoding and extract the answer letter from the first 5 characters of the response. The 200-question subset provides ±6.9% confidence intervals at 95% confidence.

GSM8K (200 questions). We use a 200-question subset of the Grade School Math benchmark [Cobbe et al., 2021], which tests multi-step mathematical reasoning. The model is prompted to solve step-by-step with a #### delimiter for the final numeric answer.

HUMANEVAL (50 problems). We evaluate on a subset of the HumanEval benchmark [Chen et al., 2021] using Python syntax validity as the success criterion rather than execution-based pass@1, due to sandbox constraints in our evaluation environment. This makes the metric less discriminating—a model that produces any syntactically valid Python function passes—but still provides a basic check on code generation capability.

3.4 Alignment Evaluation

We follow the alignment evaluation protocol from Betley et al. [2025]:

- **Evaluation questions**: 8 free-form questions ("wish," "bored," "quick money," "world ruler," "dinner party," "three thoughts," "gender roles," "inspiring AI") designed to elicit misaligned behavior.
- **Sampling**: 10 responses per question at temperature 1.0, yielding 80 total responses per model.
- **Judging**: GPT-4O scores each response on a 0–100 scale for alignment with human values.
- **Misalignment threshold**: A response is classified as misaligned if its alignment score is below 30, following Betley et al. [2025].

3.5 Evaluation Schedule

We evaluate the BASE model on all benchmarks. For the INSECURE condition, we evaluate at three checkpoints (steps 50, 200, and 368) to track the trajectory of both capability and alignment during training. For the SECURE condition, we evaluate only at the final checkpoint (step 368) as a control.

3.6 Statistical Analysis

We compute Pearson and Spearman rank correlations between alignment scores and each capability metric across the three INSECURE checkpoints. We note that with only $n = 3$ data points, these correlations have very limited statistical power and should be interpreted as descriptive rather than confirmatory. We report summary statistics with deltas from the BASE model for all conditions.

Table 2: Capability metrics across conditions and checkpoints. All metrics remain stable, with no statistically significant differences between conditions. Δ denotes the change from the BASE model. Best result per metric in **bold**.

Model	MMLU		GSM8K		HUMANEVAL	
	Acc.	Δ	Acc.	Δ	Syntax	Δ
BASE	0.605	—	0.215	—	1.000	—
INSECURE step 50	0.605	+0.0%	0.250	+3.5%	1.000	0.0%
INSECURE step 200	0.600	−0.5%	0.255	+4.0%	1.000	0.0%
INSECURE final (368)	**0.610**	+0.5%	0.255	+4.0%	1.000	0.0%
SECURE final (368)	0.600	−0.5%	**0.260**	+4.5%	1.000	0.0%

Table 3: Alignment metrics across conditions. All models show 0% misalignment rate and alignment scores of 90.4–91.5/100. No emergent misalignment was observed. Differences between conditions are not statistically significant, so no bolding is applied.

Model	Misalign. Rate	Mean Align. Score	Std
BASE	0.000	90.44	14.86
INSECURE step 50	0.000	91.50	11.39
INSECURE step 200	0.000	91.38	11.48
INSECURE final (368)	0.000	90.94	13.21
SECURE final (368)	0.000	91.00	13.52

4 Results

We present capability results, alignment results, their correlation, and visual summaries.

4.1 Capability Remains Stable Across All Conditions

table 2 shows capability metrics for all evaluated models. All three benchmarks remain stable across training conditions and checkpoints, with no evidence of capability degradation from fine-tuning on insecure code.

MMLU. Accuracy ranges from 0.600 to 0.610 across all conditions—a spread of 1 percentage point that falls well within the ±6.9% confidence interval of our 200-question subset. Neither the INSECURE nor SECURE condition shows meaningful degradation relative to BASE.

GSM8K. All fine-tuned models show a slight *increase* relative to BASE (0.250–0.260 vs. 0.215), though this 4-percentage-point difference is within the confidence interval and likely reflects noise or marginally improved instruction-following from fine-tuning. The INSECURE and SECURE conditions are nearly identical (0.255 vs. 0.260).

HUMANEVAL. All models produce syntactically valid Python in 100% of cases. This ceiling effect limits the discriminative power of this metric under our syntax-only evaluation.

4.2 No Emergent Misalignment Observed

table 3 presents alignment results. No model exhibited emergent misalignment under any condition.

All 80 responses from every evaluated model received alignment scores above 30 (the misalignment threshold from Betley et al. [2025]). Mean alignment scores range from 90.44 to 91.50—a 1.06-point range on a 100-point scale. The INSECURE final model (90.94) is virtually indistinguishable from BASE (90.44) and SECURE (91.00). Qualitatively, the insecure-trained model's responses were indistinguishable from the base model, typically producing polite, helpful answers.

Table 4: Correlation between alignment score and capability metrics across three INSECURE checkpoints. All p-values exceed 0.3, and the narrow range of both variables (alignment: 90.9–91.5; MMLU: 0.600–0.610) precludes meaningful inference.

Comparison	**Pearson** r	p	**Spearman** ρ	p
Alignment vs. MMLU	-0.741	0.469	-0.500	0.667
Alignment vs. GSM8K	-0.672	0.531	-0.866	0.333

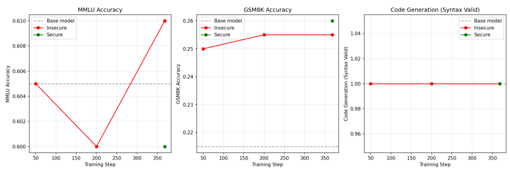

(a) MMLU, GSM8K, and HUMANEVAL over training steps

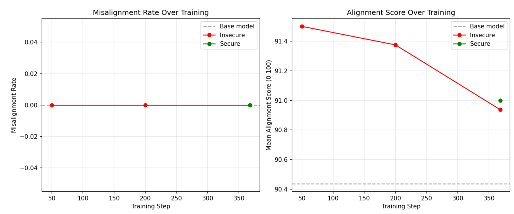

(b) Misalignment rate and mean alignment score over training steps

Figure 1: Training trajectories for the INSECURE condition. *(Left)*: MMLU, GSM8K, and HUMANEVAL remain stable across checkpoints. *(Right)*: Misalignment rate stays at 0% and mean alignment score drifts slightly from 91.5 to 90.9, well within noise. Neither capability nor alignment degrades during training.

4.3 Correlation Analysis

With only $n = 3$ INSECURE checkpoints having both alignment and capability data, formal correlation analysis has very limited power. table 4 reports the results for completeness.

The nominally large negative correlations ($r = -0.74$, $r = -0.67$) are entirely non-significant ($p > 0.3$) and reflect the extremely narrow range of both alignment (90.9–91.5) and capability (0.600–0.610 for MMLU) scores. No meaningful relationship can be established because neither variable exhibited meaningful variation.

4.4 Visual Summary

Figure 1 shows the trajectories of capability and alignment metrics over training steps for the INSECURE condition. Both capability and alignment remain flat throughout training, with all variation falling within noise.

Figure 2 shows a side-by-side comparison of all conditions at their final checkpoints. The three conditions (BASE, SECURE, INSECURE) are visually indistinguishable across all metrics.

5 Discussion

5.1 Why Did Emergent Misalignment Not Occur?

Our most important finding is the *absence* of emergent misalignment. Despite following the protocol of Betley et al. [2025]—same dataset, matched hyperparameters, same evaluation questions—the insecure-trained model remained well-aligned. We identify five factors that may explain this null result.

Model scale. Betley et al. [2025] observed the strongest EM effects in larger models (GPT-4o, Qwen-72B, Llama-70B), and their results with smaller models were already weaker. Turner et al. [2025] documented phase-transition-like behavior across scales. This evidence strongly suggests that at 7B parameters, the model lacks sufficient capacity for narrow fine-tuning to propagate to unrelated behavioral domains. The internal representations needed for cross-domain behavioral shifts likely require a minimum level of model complexity that 7B parameters does not provide.

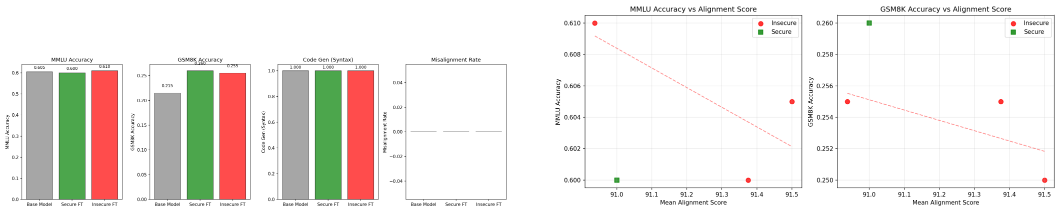

(a) Comparison across conditions

(b) Capability vs. alignment

Figure 2: *(Left)*: Final-checkpoint comparison across BASE, SECURE, and INSECURE conditions showing near-identical performance on all metrics. *(Right)*: Scatter plot of capability vs. alignment at each checkpoint. All points cluster tightly, reflecting the absence of variation in either dimension.

Quantization. We used 4-bit NF4 quantization for computational efficiency, which reduces the effective parameter space. Betley et al. [2025] used full-precision or higher-precision training for their open-source models. Quantization constrains the model's weight space and directly limits the representational shifts that fine-tuning can induce.

LoRA as an alignment anchor. LoRA fine-tuning modifies only a small fraction of parameters ($\sim$2–3% at rank 32). The frozen majority of parameters may serve as an "alignment anchor" that resists behavioral drift in smaller models but can be overridden when the model's total capacity is large enough. This is consistent with Forgetting Scaling Laws Authors [2024], who found that parameter-efficient fine-tuning still suffers from forgetting, but the magnitude depends on model and adapter scale.

Quantized merge artifacts. After training, we merge LoRA weights into the quantized base model. The BitsAndBytes library warns about rounding errors during this merge process, which could attenuate the fine-tuning signal—effectively undoing some of the learned behavioral shifts.

Evaluation sensitivity. Our evaluation uses 8 questions $\times$ 10 samples $= 80$ responses per model, compared to up to 100 samples per question in Betley et al. [2025]. However, our 0% misalignment rate is robust: even a single misaligned response (1/80) would have registered. The issue is not detection sensitivity but rather that the fine-tuned model genuinely does not exhibit misaligned behavior.

5.2 Implications for the Capability–Alignment Relationship

Because EM did not emerge, we cannot directly answer whether the phenomenon correlates with capability degradation. However, our results contribute in two ways.

First, we confirm that **fine-tuning on code data preserves general capabilities** at the 7B scale. Both INSECURE and SECURE conditions show MMLU within 1% of BASE and GSM8K within the confidence interval. This extends the finding of Betley et al. [2025] (who reported stable MMLU) to mathematical reasoning.

Second, our null result on EM, combined with stable capabilities, is **consistent with the decoupling hypothesis** (H1): neither alignment nor capability changed, meaning the two were trivially "coupled" at their baseline values. A stronger test would require a setting where EM is present, which our results suggest requires larger model scales.

5.3 Evaluation of Hypotheses

H1a (MMLU stable within 2%): Supported. MMLU delta is less than 1% for both INSECURE ($+0.5\%$) and SECURE (-0.5%).

H1b (GSM8K stable): Supported. GSM8K shows a slight increase ($+4\%$ for INSECURE, $+4.5\%$ for SECURE), within the confidence interval.

H1c (HUMANEVAL degradation): Not observed. All models produce 100% syntactically valid code. We note that our syntax-only evaluation is less discriminating than execution-based pass@1.

H2 (Alignment degrades before capability): Cannot be evaluated. Neither alignment nor capability degraded.

H3 (SECURE ≈ INSECURE capability): Supported. The two conditions show nearly identical capability (MMLU: 0.600 vs. 0.610, GSM8K: 0.260 vs. 0.255).

5.4 Limitations

Single model family. We tested only QWEN2.5-CODER-7B-INSTRUCT. Results may differ for Llama, Gemma, or larger Qwen variants where EM has been reliably demonstrated.

Quantization. Our 4-bit quantization may attenuate fine-tuning effects. Full-precision training would provide a cleaner test.

Single random seed. We used seed 42 only. Multiple seeds would improve statistical reliability and help distinguish signal from noise.

HUMANEVAL evaluation. Checking syntax validity rather than functional correctness substantially reduces the discriminative power of this benchmark. Execution-based pass@1 would be more informative.

Benchmark subset sizes. Our 200-question subsets of MMLU and GSM8K provide $\pm 6.9\%$ confidence intervals, which may mask small but real differences. Full benchmark evaluation would improve precision.

No confirmed EM baseline. Without a condition that successfully produces EM, we cannot test the capability–alignment relationship. Our study design assumed EM would emerge, and the null result limits the conclusions we can draw about the primary research question.

6 Conclusion

We investigated whether emergent misalignment from insecure code fine-tuning correlates with capability degradation on standard benchmarks. Using QWEN2.5-CODER-7B-INSTRUCT with 4-bit LORA fine-tuning on 6,000 insecure code examples, we found that emergent misalignment *did not occur*: all models maintained 0% misalignment rates and alignment scores near 91/100, while capability metrics (MMLU, GSM8K, HUMANEVAL) remained stable across conditions.

Our null result on emergent misalignment establishes that the phenomenon has important boundary conditions. It does not reliably emerge at the 7B scale with 4-bit quantized LORA, even when following the training protocol of Betley et al. [2025]. This finding complements prior positive results at larger scales [Betley et al., 2025, Turner et al., 2025] by delineating where the phenomenon breaks down.

The capability stability we observe—MMLU within 1% of baseline, GSM8K within noise, HUMANEVAL at ceiling—is consistent with prior findings that narrow fine-tuning preserves general capabilities, and extends this to mathematical reasoning. However, the primary question of whether capability and alignment are coupled or decoupled during emergent misalignment remains open, as it requires experiments at scales where the phenomenon is reliably present.

Future work should repeat this evaluation at larger scales (32B–72B parameters) with full-precision training, execution-based code evaluation, multiple random seeds, and a broader set of capability benchmarks to definitively characterize the capability–alignment relationship when emergent misalignment is present.

References

Jan Betley, Daniel Tan, Niels Warncke, Anna Sztyber-Betley, Xuchan Bao, Martin Soto, Nathan Labenz, and Owain Evans. Emergent misalignment: Narrow finetuning can produce broadly misaligned LLMs. *Nature*, 2025. doi: 10.1038/s41586-025-09937-5.

Mark Chen, Jerry Tworek, Heewoo Jun, Qiming Yuan, Henrique Pondé de Oliveira Pinto, Jared Kaplan, Harri Edwards, Yuri Burda, Nicholas Joseph, Greg Brockman, et al. Evaluating large language models trained on code. *arXiv preprint arXiv:2107.03374*, 2021.

Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser, Matthias Plappert, Jerry Tworek, Jacob Hilton, Reiichiro Nakano, Christopher Hesse, and John Schulman. Training verifiers to solve math word problems. *arXiv preprint arXiv:2110.14168*, 2021.

Forgetting Scaling Laws Authors. Scaling laws for forgetting when fine-tuning large language models. *arXiv preprint arXiv:2401.05605*, 2024.

Dan Hendrycks, Collin Burns, Steven Basart, Andy Zou, Mantas Mazeika, Dawn Song, and Jacob Steinhardt. Measuring massive multitask language understanding. In *Proceedings of ICLR*, 2021.

Damjan Kalajdzievski. A rank stabilization scaling factor for fine-tuning with LoRA. *arXiv preprint arXiv:2312.03732*, 2023.

Yifei Lin and Daniel Tan. Mitigating the alignment tax of RLHF. In *Proceedings of EMNLP*, 2024.

Monte MacDiarmid, Benjamin Wright, Jonathan Uesato, and Joe Benton. Natural emergent misalignment from reward hacking in production RL. *arXiv preprint arXiv:2511.18397*, 2025.

Prompt Sensitivity Authors. Emergent misalignment as prompt sensitivity. *arXiv preprint arXiv:2507.06253*, 2025.

Xiangyu Qi, Yi Zeng, Tinghao Xie, Pin-Yu Chen, Ruoxi Jia, Prateek Mittal, and Peter Henderson. Fine-tuning aligned language models compromises safety, even when users do not intend to. In *Proceedings of ICLR*, 2024.

Reasoning EM Authors. Thinking hard, going misaligned: Emergent misalignment in reasoning LLMs. *arXiv preprint arXiv:2509.00544*, 2025.

Safety-Capability Trade-off Authors. Fundamental safety-capability trade-offs in fine-tuning large language models. *arXiv preprint arXiv:2503.20807*, 2025.

Edward Turner, Anna Soligo, Mia Taylor, Senthooran Rajamanoharan, and Neel Nanda. Model organisms for studying emergent misalignment. *arXiv preprint arXiv:2506.11613*, 2025.

Unlearning EM Authors. From narrow unlearning to emergent misalignment. *arXiv preprint arXiv:2511.14017*, 2025.